



Cloud & DevOps Best Practices

IGNITE BRILLIANCE

Contents

Cloud & DevOps Standards	3
1. Why This Document Exists.....	3
2. Who Should Use This?	3
3. Architecture Principles	3
4. Scalability & Performance.....	4
5. Security (Built In, Not Bolted On)	4
6. CI/CD & Deployment	5
7. Monitoring, Alerts, and Observability	5
8. Logging & Auditability	6
9. Backups & Disaster Recovery.....	6
10. Governance & Compliance	6
11. Change & Incident Management	7
12. Knowledge Sharing & Training.....	7
13. Review & Maintenance	7

Cloud & DevOps Standards

A Practical Guide to Building Resilient, Scalable, and Secure Systems

1. Why This Document Exists

This guide outlines how we approach cloud infrastructure and DevOps—not just to tick boxes, but to make sure our systems are **resilient, secure, scalable**, and ready to support real business outcomes. It's grounded in best practices from **CMMI, ISO/IEC 27001, 20000, and 12207**, but more importantly, it's shaped by what actually works for modern engineering teams.

2. Who Should Use This?

If you're:

- Designing systems or infrastructure
- Building and deploying code
- Managing environments
- Monitoring performance or responding to incidents
- Auditing for compliance or security

...then this document is for you.

3. Architecture Principles

We build systems to be:

- **Cloud-native:** Microservices, containers, serverless where it makes sense.
- **Resilient by design:** We assume failure will happen—and design with retries, fallbacks, and timeouts.
- **Highly available:** Across multiple zones or regions, with no single point of failure.
- **Infrastructure-as-Code:** If it's not in version control, it doesn't exist.

Best practices:

- Use Infrastructure as Code (IaC) to deploy infrastructure, example:
 - Cloud Development Kit (CDK) for AWS deployments.
 - Terraform or Pulumi for cloud agnostic deployments
- Make services stateless wherever possible.

- Include architecture diagrams in your project repo (and keep them updated!).
 - Use managed services when possible
-

4. Scalability & Performance

Scalability isn't optional—our systems should handle growth without breaking a sweat.

What we do:

- Auto-scaling groups and serverless wherever it fits (e.g., AWS Lambda, GCP Cloud Run).
- Load testing with tools like **k6**, **JMeter**, or **Artillery** before go-live.
- Use of **caching layers** (Redis, CDN, browser caching) to offload repeated work.

Design tips:

- Plan for growth (3x traffic) in your architecture.
 - Watch memory leaks or thread bottlenecks in long-running services.
 - Run performance regression tests in CI.
-

5. Security (Built In, Not Bolted On)

Security is everyone's job, but especially ours.

What we enforce:

- **Least privilege access** using IAM roles/policies.
- **Secrets management** using AWS Secrets Manager, HashiCorp Vault, or environment-specific key stores.
- TLS 1.2+ everywhere.
- Container scanning with tools like Trivy or Snyk.
- Mandatory static code analysis (e.g., SonarQube, CodeQL) for every pull request.

Regular practices:

- Threat modeling for new features.
 - Security retros after incidents.
 - Quarterly penetration testing.
-

6. CI/CD & Deployment

We aim for fast, safe, and reversible releases.

Pipeline guidelines:

- Git-based pipelines (Bitbucket Pipelines).
- Every commit triggers tests, builds, security scans (SonarQube, Trivy).
- Tag all releases; no “Thursday deploys” without business approval.

Deployment patterns:

- **Blue/green, canary, or rolling** updates depending on risk level.
- Use of **feature flags** for controlled rollouts.
- Pre-prod mirror environments for testing.

Always keep **rollback instructions** ready in your release notes.

7. Monitoring, Alerts, and Observability

“If it’s not monitored, it doesn’t exist.”

What we monitor:

- **Service uptime** (HTTP 2xx/5xx ratios, latency, error rates).
- **Infrastructure health** (CPU, memory, disk I/O).
- **Business metrics** (signups, revenue-impacting flows).

Tools we love:

- Prometheus + Grafana
- Cloud-native tools (CloudWatch, Azure Monitor, GCP Operations)
- Real-time alerting via Teams or PagerDuty

Every alert must:

- Be actionable (no “noise” alerts)
 - Link to a **runbook**
 - Include a severity level and clear ownership
-

8. Logging & Auditability

Logs help us debug, understand, and stay compliant.

Logging rules:

- Use structured logging (e.g., JSON format).
- Mask PII and secrets in logs—ALWAYS.
- Ship logs to a centralized location (ELK Stack, CloudWatch, or equivalent).
- Retain logs as per company policy (typically 90 days for app logs, 1 year for audit logs).

Audit trails:

- Required for privileged access and config changes.
 - Automatically captured in cloud provider dashboards or SIEM tools.
-

9. Backups & Disaster Recovery

Stuff breaks. Be ready.

Backup practices:

- Daily backups for DBs, weekly for file systems.
- Encrypted backups stored in separate region or zone.
- Quarterly restore drills to test RPO/RTO.

DR documentation must include:

- Contacts
 - Systems impacted
 - Restoration priority
 - Checklist for DR testing
-

10. Governance & Compliance

This is where ISO and CMMI meet the real world.

Core practices:

- Tag all cloud resources by environment, team, project and client.
- Automate policy enforcement with tools like **AWS Config**.

- Use CMDBs or inventory tools (like ServiceNow).
 - Maintain evidence for regular audits (change logs, access logs, vulnerability scans).
-

11. Change & Incident Management

All infra and app changes:

- Go through Git-based PR process with peer reviews.
- Are traceable (commit → ticket → release).
- Include a rollback plan.

Incidents:

- Logged in your ITSM tool (e.g., Jira, ServiceNow).
 - Follow our incident playbook and severity matrix.
 - Include postmortems with root cause analysis and next steps.
-

12. Knowledge Sharing & Training

We write things down. We teach what we know.

Expect to:

- Maintain internal docs in Confluence or Loop or SharePoint.
- Include architecture, runbooks, and "how to" guides in every repo.
- Host regular internal knowledge-sharing sessions.

For continuous improvement:

- Track **DORA Metrics** (Deployment Frequency, Lead Time, MTTR, Change Failure Rate).
 - Run retros after big launches or outages.
 - Encourage feedback and improve documentation regularly.
-

13. Review & Maintenance

This document isn't static. It evolves as we do.

Review schedule:

- Every 6 months or after a major incident or audit.
 - Managed by the Engineering Leadership Team.
 - Input from all job families not just developers.
-